



Word2Vector

Speaker : 蘇佳益

Advisor : 陳聰毅

國立高雄科技大學建工校區電子工程系

[https://github.com/chiayisu/Artificial Intelligence Course](https://github.com/chiayisu/Artificial_Intelligence_Course)

Agenda

- Introduction
- N-Grams
- Introduction to Word Vector
- Complexity of Word Vector
- Negative Sampling
- Hierarchical Softmax
- Practical Topics
- Results
- References

Introduction

- N-Grams
 - Robustness
 - Simplicity
 - However: we cannot know similarity between words.
- WordVector (Mikolov et al.)
 - Represent words as dense vector
 - Able to perform simple algebraic operation
 - E.g.: $\text{Vector}(\text{queen}) = \text{Vector}(\text{king}) - \text{Vector}(\text{men}) + \text{Vector}(\text{women})$



N-Grams

N-Grams: Introduction

- The simplest model to assign probabilities to each sentence
- Sequence of N words
- N can be any positive number. However: Usually (4-10)
- Our example will be $N=2$ (Bi-Gram)

N-Grams: Introduction

- Suppose we have a sentence, it is shown that, the word w is “that” given the history h “it is shown”. Let’s compute $p(w|h)$.
- $p(w|h) = p(\textit{that}|\textit{it is shown}) = \frac{c(\textit{it is shown that})}{c(\textit{it is shown})}$

The computation above is problematic.

N-Grams: Why is it problematic?

- As we mentioned in “Introduction Course”, Natural Language often
 - evolves.
 - sometimes doesn't follow grammar.
 - be equivocal.
- We may not find the sentences of the history in our database or web.

Thus, We need to consider few words

N-Grams: Bi-Grams

- According to Markov assumption
 - we only need to take few words into consideration.
 - Therefore: in previous example, we only need to calculate $p(\text{that} \mid \text{shown})$ in bi-grams model.
- Markov model is that we can predict the future words without looking too far into the past.

N-Grams: Bi-Grams - Estimation

- We estimate the probability intuitively by using Maximum Likelihood Estimation. The formula is as follow.
- $$p(w_n|w_{n-1}) = \frac{C(W_{n-1}W_n)}{\sum_w C(W_{n-1}W)} = \frac{C(W_{n-1}W_n)}{C(W_{n-1})}$$

Let's walk through an example.

N-Grams: Example

- Take a look at the following sentences
 - I am Ian
 - Ian am I
 - My name is Ian
- The following probabilities are the bi-grams probability in this corpus.
- $P(am | I) = \frac{1}{2}$
- $P(Ian | am) = \frac{1}{2}$
- $P(am | Ian) = \frac{1}{3}$

Apparently, N-grams still make no sense about similarity.



Introduction to Word Vector

Previous Work

- Downside of Previous Model
 - computational costly
- E.g.: NNLM contains
 - feedforward neural network with a linear projection layer
 - non-linear hidden layer

Word2Vector Model (Mikolov et al.)

- Bag-of-Words model (CBOW)
- Continuous Skip-gram model



Complexity of Word Vector

Complexity Notation

- The following model's complexity will be represented as follow.

➤ $O = E * T * Q$

- where E is the number of training epoch, usually 3 - 50, T is the number of words, usually 1billion and Q will be defined depending on the further model.

Complexity of NNLM

- NNLM is a language model that consists of input, projection, hidden and output layer. Given the N previous words, it will predict the next word.
- Complexity of Projection Layer
 - $N * D$, where N is the number of previous words and D is its dimension.
- Complexity of Hidden Layer
 - $N * D * H$, where $N * D$ is from input layer and H is from Hidden layer.
- Complexity of Output Layer
 - $H * V$, where H is the number of hidden layer and V is the size of all vocabulary.
- Thus, its complexity, in total, is $Q = N * D + N * D * H + H * V$

CBOW

- CBOW Model
 - Is similar to NNLM
 - Predicts future words based on given history words
 - Is also similar to bag-of words model
 - However: heaviest layer hidden layer in NNLM is removed
 - Therefore: Its projection layer is shared with all words
- Complexity of CBOW
 - $Q = N * D + D * V$

Skip-Gram

- Skip-Gram model
 - Is opposite from the CBOW.
 - Predicts the context based on given word
- Complexity of Skip-Gram
 - $Q = C * (D + D * V)$, where C is the distance of the word. The size of C is double because it consists of size from history and future.

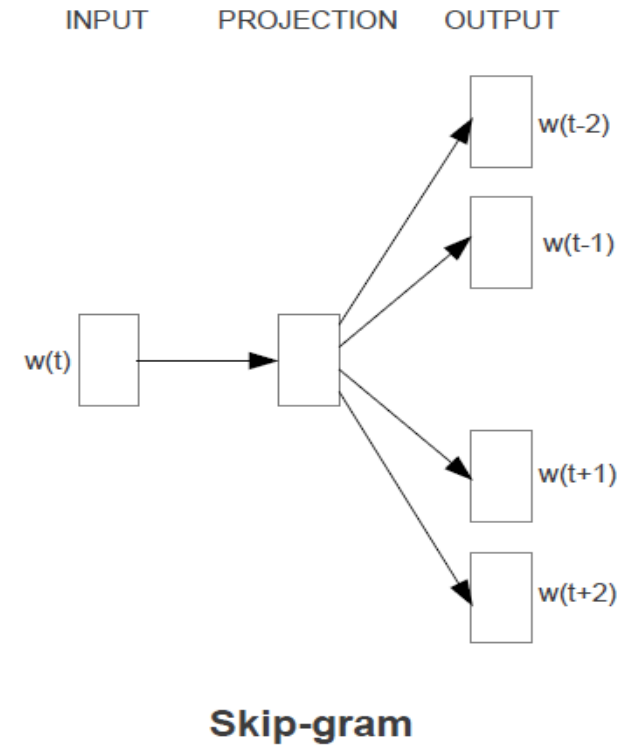
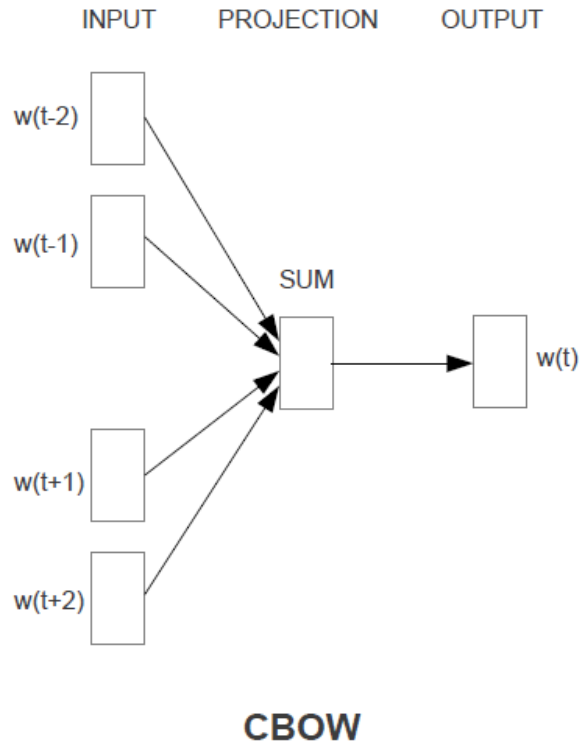
Complexity: Binary Tree Representation

- Complexity of Binary Tree
 - $\log(n)$
- Therefore: the complexity of representing vocabulary as binary tree is $\log(v)$
- However: hierarchical Softmax is required instead of Softmax.



Algorithm of CBOW and Skip-Gram

Architecture of CBOW and Skip-Gram





CBOW

Notation for Algorithm

- w_i : word i from vocabulary V
- $\omega \in R^{n*|V|}$: input word matrix
- ω_i : i – th coulumn of ω
- $U \in R^{|V|*n}$: ouput word matrix
- u_i : i – th coulumn of U

Algorithm of CBOW

1. Generate one hot word vector for window size m :
 $(x^{(c-m)}, \dots, x^{(c-1)}, \dots, x^{(c+m)} \in R^{|V|})$
2. Get embedded word vector ($\hat{v} = \omega x^{(c-m)} + \dots + \omega x^{(c+m)} \in R^n$)
3. Average it to get $\hat{v} = \frac{\hat{v}}{2m} \in R^n$
4. Generate score vector $z = U * \hat{v} \in R^{|V|}$
5. Turn the score into probability $y = \text{softmax}(z) \in R^{|V|}$
6. Run gradient descent to update input and output word matrix

Update Input and Output Word Matrix

- Information theory is employed when
 - we want to learn the probability from true probability.
- Cross-Entropy loss is derived as follows.
- $H(\hat{y}, y) = -y_i \log(\hat{y}_i)$
- Take a look at example, if $\hat{y}_i = 1$, our H will be 0, which is no loss. If $\hat{y}_i = 0.01$, our H will be ≈ 4.605

Update Input and Output Word Matrix (CONT.)

- Thus, we need to minimize our cross-entropy.
- Let $J = -\log P(w_c | w_{c-m}, \dots, w_{c+m})$

$$= -\log P(u_c | \hat{v})$$

$$= -\log \frac{\exp(u_c^T \hat{v})}{\sum_{j=1}^{|V|} \exp(u_j^T \hat{v})}$$

$$= -u_c^T \hat{v} + \log \left(\sum_{j=1}^{|V|} \exp(u_j^T \hat{v}) \right)$$

Gradient descent is required to update \hat{v} and u_c



Skip-Gram

Algorithm of Skip-Gram

- Generate one hot input vector $x \in R^{|V|}$ of center word
- Get embedded word vector for center word $v_c = \omega x \in R^n$
- Generate a score vector $z = U * v_c$
- Turn score vector into probability $\hat{y} = \text{softmax}(z)$
- Update input and output word matrix

Algorithm of Skip-Gram (cont.)

- Basically, it's similar to CBOW. Input and Output matrices are required to update
- However: Skip-Gram assumes output words are completely independent as Naïve Bayes algorithm.

- Let $J = -\log P(w_{c-m}, \dots, w_{c+m} | w_c)$
$$\begin{aligned} &= -\log\left(\prod_{j=0, j \neq m}^{2m} P(w_{c-m+j} | w_c)\right) \\ &= -\log\left(\prod_{j=0, j \neq m}^{2m} P(u_{c-m+j} | v_c)\right) \\ &= -\log\left(\prod_{j=0, j \neq m}^{2m} \frac{\exp(u_{c-m+j}^T v_c)}{\sum_{k=1}^{|V|} \exp(u_k^T v_c)}\right) \\ &= -\sum_{j=0, j \neq m}^{2m} u_{c-m+j}^T v_c + 2m \log\left(\sum_{k=1}^{|V|} \exp(u_k^T v_c)\right) \end{aligned}$$



Negative Sampling

Introduction

- Disadvantage of Softmax
 - Computational Costly
 - Reason: takes $O(|V|)$ time
- Negative Sampling
 - Only needs to loop over noise distribution
 - Is based on skip-gram model

Objective Function

- Let's denote (w, c) to word and context, $P(D = 1|w, c)$ to the probability that (w, c) comes from corpus data, and $P(D = 0|w, c)$ vice versa. Therefore, we can model $P(D = 1|w, c)$ with sigmoid function as follow.
- $$P(D = 1|w, c, \theta) = \sigma(v_c^T v_w) = \frac{1}{1 + e^{-(v_c^T v_w)}}$$

Objective Function (CONT.)

- New objective function is built to maximize the probability of a word and context in the corpus data. It's derived bellow. We also denote θ as the parameters of input and output matrix.

$$\begin{aligned} \bullet \theta &= \operatorname{argmax}_{\theta} \prod_{(w,c) \in D} p(D = 1|w, c, \theta) \prod_{(w,c) \in \hat{D}} p(D = 0|w, c, \theta) \\ &= \operatorname{argmax}_{\theta} \prod_{(w,c) \in D} p(D = 1|w, c, \theta) \prod_{(w,c) \in \hat{D}} (1 - p(D = 1|w, c, \theta)) \\ &= \operatorname{argmax}_{\theta} \sum_{(w,c) \in D} \log(p(D = 1|w, c, \theta)) + \sum_{(w,c) \in \hat{D}} \log(1 - p(D = 1|w, c, \theta)) \\ &= \operatorname{argmax}_{\theta} \sum_{(w,c) \in D} \log\left(\frac{1}{1 + \exp(-u_w^T v_c)}\right) + \sum_{(w,c) \in \hat{D}} \log\left(1 - \frac{1}{1 + \exp(-u_w^T v_c)}\right) \\ &= \operatorname{argmax}_{\theta} \sum_{(w,c) \in D} \log\left(\frac{1}{1 + \exp(-u_w^T v_c)}\right) + \sum_{(w,c) \in \hat{D}} \log\left(\frac{1}{1 + \exp(u_w^T v_c)}\right) \end{aligned}$$

Objective Function (CONT.)

- Because maximizing log likelihood is the same as minimizing negative log likelihood, our function will be revised below.

- $$J = -\sum_{(w,c) \in D} \log\left(\frac{1}{1+\exp(-u_w^T v_c)}\right) - \sum_{(w,c) \in \hat{D}} \log\left(\frac{1}{1+\exp(u_w^T v_c)}\right)$$

\hat{D} is a false corpus, where unnatural sentences contain in this corpus. \hat{D} can be generated by randomly sampling from word bank.

Objective Function (CONT.)

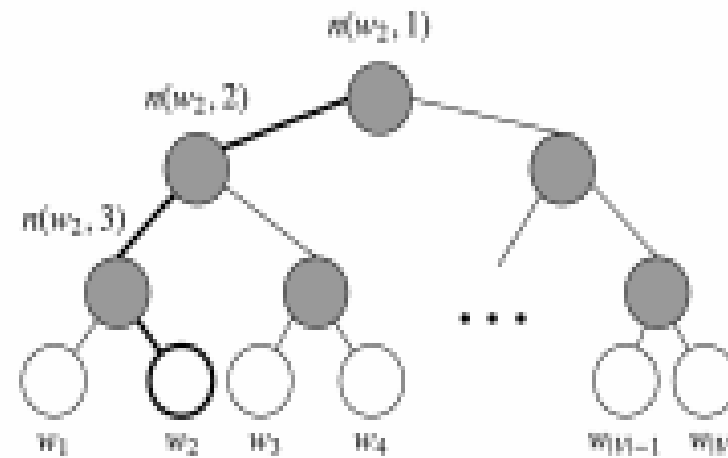
- New objective function for skip-gram model is, thus, as follows
- $-\log(\sigma(u_{c-m+j}^T * v_c)) - \sum_{k=1}^K \log(\sigma(-\hat{u}_k^T * v_c))$
- New objective function for CBOW model is, thus, as follows
- $-\log(\sigma(u_c^T * v_c)) - \sum_{k=1}^K \log(\sigma(-\hat{u}_k^T * v_c))$
- $\{\hat{u}_k | k = 1 \dots K\}$ are sampled from $P(w)$. $P(W)$ is generally unigram model raise to $\frac{3}{4}$ because the unnatural words are more likely to be sampled.



Hierarchical Softmax

Introduction

- Hierarchical softmax
 - represents all words as binary tree.
 - Each leaf is a word and each node is a vector that model is going to learn.
- Advantage of using hierarchical softmax
 - Only cost $O(\log(V))$
- Disadvantage of using hierarchical softmax
 - Cannot be parallelized



Notation

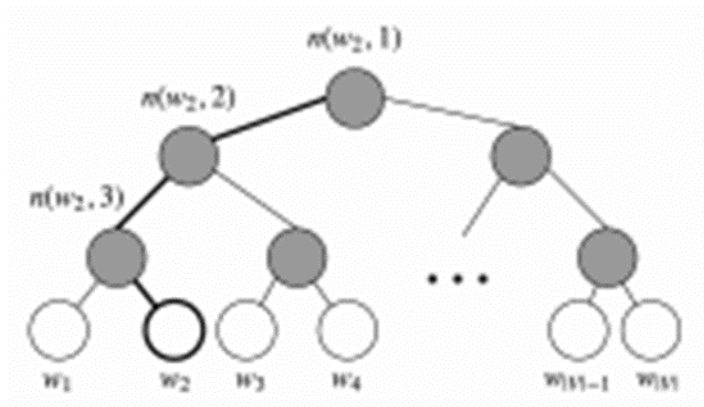
- $L(w)$: number of node from root to leaf w
- $n(w,i)$: i -th node on the path
- $ch(i)$: inner node n 's children

Probability Calculation

- The probability is, then, as follow.
- $P(w|w_i) = \prod_{j=1}^{L(w)} \sigma([n(w, j + 1) = ch(n(w, j))] * v_{n(w, j)}^T v_{w_i})$
- where $[x] = \begin{cases} 1 & \text{if } x \text{ is true} \\ -1 & \text{otherwise} \end{cases}$
- We also assume that $ch(n)$ is the left node of n . if path goes left, $[n(w, j + 1) =$

Example

- take the graph bellow for example. The approach to calculate $p(w_2|w_i)$ is as follow.
- $$p(w_2|w_i) = p(n(w_2, 1), left) * p(n(w_2, 2), left) * p(n(w_2, 3), right)$$
$$= \sigma(v_{n(w_2,1)}^T v_{w_i}) * \sigma(v_{n(w_2,2)}^T v_{w_i}) * \sigma(-v_{n(w_2,3)}^T v_{w_i})$$



Our goal is the same as before,
which is minimize $-\log p(w|w_i)$



Practical Topics

Subsampling of Frequent Words

- Meaningless words
 - For instance: a, an, the
 - are often in the large corpora
 - Appear most of time
- However: Skip-gram model doesn't benefit from this kind of words
- Therefore: we need to counter the imbalance.
- $P(w_i) = 1 - \sqrt{\frac{t}{f(w_i)}}$, $f(w_i)$ is the frequency of the word and is a chosen threshold. It is typically around 10^{-5} .

Learning Phrase

- Some words appear together as phrases
 - E.g. New York Times
- This paper (Mikolov et al.) only uses simple model based on unigram and bigram counts.
- $score(w_i, w_j) = \frac{count(w_i, w_j) - \delta}{count(w_i) * count(w_j)}$
- δ is a coefficient to avoid too many infrequent phrases. We choose the phrases that have the score above the threshold.
- This algorithm is often run 2-4 passes with declining threshold.



Results

Examples of Semantic and Syntactic Test

| Type of relationship | Word Pair 1 | | Word Pair 2 | |
|-----------------------|-------------|------------|-------------|---------------|
| Common capital city | Athens | Greece | Oslo | Norway |
| All capital cities | Astana | Kazakhstan | Harare | Zimbabwe |
| Currency | Angola | kwanza | Iran | rial |
| City-in-state | Chicago | Illinois | Stockton | California |
| Man-Woman | brother | sister | grandson | granddaughter |
| Adjective to adverb | apparent | apparently | rapid | rapidly |
| Opposite | possibly | impossibly | ethical | unethical |
| Comparative | great | greater | tough | tougher |
| Superlative | easy | easiest | lucky | luckiest |
| Present Participle | think | thinking | read | reading |
| Nationality adjective | Switzerland | Swiss | Cambodia | Cambodian |
| Past tense | walking | walked | swimming | swam |
| Plural nouns | mouse | mice | dollar | dollars |
| Plural verbs | work | works | speak | speaks |

Accuracy on Semantic-Syntactic Word Relationship Test Set with CBOW

| Dimensionality / Training words | 24M | 49M | 98M | 196M | 391M | 783M |
|---------------------------------|------|------|------|------|------|------|
| 50 | 13.4 | 15.7 | 18.6 | 19.1 | 22.5 | 23.2 |
| 100 | 19.4 | 23.1 | 27.8 | 28.7 | 33.4 | 32.2 |
| 300 | 23.2 | 29.2 | 35.3 | 38.6 | 43.7 | 45.9 |
| 600 | 24.0 | 30.1 | 36.5 | 40.8 | 46.6 | 50.4 |

- The much data and higher dimension we use to train model, The higher accuracy we get.

Comparison of Model Trained on Same Data, with 640-Dimensional Word Vectors

| Model Architecture | Semantic-Syntactic Word Relationship test set | |
|-----------------------|---|------------------------|
| | Semantic Accuracy [%] | Syntactic Accuracy [%] |
| RNNLM | 9 | 36 |
| NNLM | 23 | 53 |
| CBOW | 24 | 64 |
| Skip-gram | 55 | 59 |

Comparison of Models Using Distbelied Distributed Framework

| Model | Vector Dimensionality | Training words | Accuracy [%] | | | Training time [days x CPU cores] |
|-----------|-----------------------|----------------|--------------|-----------|-------|----------------------------------|
| | | | Semantic | Syntactic | Total | |
| NNLM | 100 | 6B | 34.2 | 64.5 | 50.8 | 14 x 180 |
| CBOW | 1000 | 6B | 57.3 | 68.9 | 63.7 | 2 x 140 |
| Skip-gram | 1000 | 6B | 66.1 | 65.1 | 65.6 | 2.5 x 125 |

- It takes too long to finishing training 1000 dimension on NNLM model.

Example of Word Pair Relationships Using Best Skip-Gram Model

| Relationship | Example 1 | Example 2 | Example 3 |
|----------------------|---------------------|-------------------|----------------------|
| France - Paris | Italy: Rome | Japan: Tokyo | Florida: Tallahassee |
| big - bigger | small: larger | cold: colder | quick: quicker |
| Miami - Florida | Baltimore: Maryland | Dallas: Texas | Kona: Hawaii |
| Einstein - scientist | Messi: midfielder | Mozart: violinist | Picasso: painter |
| Sarkozy - France | Berlusconi: Italy | Merkel: Germany | Koizumi: Japan |
| copper - Cu | zinc: Zn | gold: Au | uranium: plutonium |
| Berlusconi - Silvio | Sarkozy: Nicolas | Putin: Medvedev | Obama: Barack |
| Microsoft - Windows | Google: Android | IBM: Linux | Apple: iPhone |
| Microsoft - Ballmer | Google: Yahoo | IBM: McNealy | Apple: Jobs |
| Japan - sushi | Germany: bratwurst | France: tapas | USA: pizza |

- According to paper (Tomas Mikolov et al.), it only reaches 60% on accuracy, but it will be improved if we train on larger data sets.

Accuracy of Various 300-Dimensional Skip-Gram Model

| Method | Time [min] | Syntactic [%] | Semantic [%] | Total accuracy [%] |
|---|------------|---------------|--------------|--------------------|
| NEG-5 | 38 | 63 | 54 | 59 |
| NEG-15 | 97 | 63 | 58 | 61 |
| HS-Huffman | 41 | 53 | 40 | 47 |
| NCE-5 | 38 | 60 | 45 | 53 |
| The following results use 10^{-5} subsampling | | | | |
| NEG-5 | 14 | 61 | 58 | 60 |
| NEG-15 | 36 | 61 | 61 | 61 |
| HS-Huffman | 21 | 52 | 59 | 55 |

- NEG-k means Negative Sampling with k negative samples.

Accuracy of Skip-Gram Models on the Phrase Analogy Dataset

| Method | Dimensionality | No subsampling [%] | 10^{-5} subsampling [%] |
|------------|----------------|--------------------|---------------------------|
| NEG-5 | 300 | 24 | 27 |
| NEG-15 | 300 | 27 | 42 |
| HS-Huffman | 300 | 19 | 47 |

| | NEG-15 with 10^{-5} subsampling | HS with 10^{-5} subsampling |
|---------------|-----------------------------------|-------------------------------|
| Vasco de Gama | Lingsugur | Italian explorer |
| Lake Baikal | Great Rift Valley | Aral Sea |
| Alan Bean | Rebecca Naomi | moonwalker |
| Ionian Sea | Ruegen | Ionian Islands |
| chess master | chess grandmaster | Garry Kasparov |

- we got the lowest accuracy on HS-Huffman without subsampling, but we reached highest accuracy when we trained with subsampling.
- This result reached 72% of accuracy.

Comparison of Model with Previous Models

| Model (training time) | Redmond | Havel | ninjutsu | graffiti | capitulate |
|-------------------------------|--|---|--|------------------------------------|---|
| Collobert (50d) (2 months) | conyers lubbock keene | plauen dzerzhinsky osterreich | reiki kohona karate | cheesecake gossip dioramas | abdicate accede rearm |
| Turian (200d) (few weeks) | McCarthy Alston Cousins | Jewell Arzu Ovitz | - - - | gunfire emotion impunity | - - - |
| Mnih (100d) (7 days) | Podhurst Harlang Agarwal | Pontiff Pinochet Rodionov | - - - | anaesthetics monkeys Jews | Mavericks planning hesitated |
| Skip-Phrase (1000d, 1 day) | Redmond Wash. Redmond Washington Microsoft | Vaclav Havel president Vaclav Havel Velvet Revolution | ninja martial arts swordsmanship | spray paint grafitti taggers | capitulation capitulated capitulating |

- Empty means words was not in the vocabulary sets.

References

- Mikolov et al., Efficient Estimation of Word Representations in Vector Space
- Mikolov et al., Distributed Representation of Words and Phrases and their Compositionality
- Manning et al., CS224N WordVector Lecture Note in Stanford University
- Alex Minnaar Word2Vec Tutorial Part II: The Continuous Bag-of-Words Model
- Jurafsky & Martin, Speech and Language Processing N-grams